Device (Orin-class 4B/8B) vs cloud (H100 32B) — accuracy per workload
upper block n = 845 / 400 / 401; lower block = the Orin's 120-question S-EMBER sample only

cell (model · construction · backend · hardware)	n	acc %	Wilson 95%
S-EMBER device 4B-I (A6000 llama.cpp F16, DEV-RECENT)	845	31.5	[28.4, 34.7]
S-EMBER device 4B-T (A6000 llama.cpp F16, DEV-RECENT)	845	33.4	[30.3, 36.6]
S-EMBER cloud 32B-I RICH, K1 (H100 vLLM)	845	39.5	[36.3, 42.9]
S-EMBER cloud 32B-T RICH, K2 (H100 vLLM)	845	33.8	[30.7, 37.1]
ERQA device 4B-I (A6000 HF sdpa)	400	42.0	[37.3, 46.9]
ERQA device 4B-T (A6000 HF sdpa)	400	45.0	[40.2, 49.9]
ERQA device 8B-I, StudyQ arm D (A6000 HF sdpa)	400	42.5	[37.7, 47.4]
ERQA device 8B-T, StudyQ arm C (A6000 HF sdpa)	400	43.5	[38.7, 48.4]
ERQA cloud 32B-I, StudyR arm C (H100 HF fa2)	400	46.8	[41.9, 51.6]
ERQA cloud 32B-T, StudyR arm D (H100 HF fa2)	400	48.8	[43.9, 53.6]
MV-RoboBench device 4B-I (A6000 HF sdpa)	401	36.9	[32.3, 41.7]
MV-RoboBench device 4B-T (A6000 HF sdpa)	401	43.6	[38.9, 48.5]
MV-RoboBench cloud 32B-I, StudyS2 arm C (H100 HF fa2)	401	48.1	[43.3, 53.0]
MV-RoboBench cloud 32B-T, StudyS2 arm D (H100 HF fa2)	401	43.9	[39.1, 48.8]
— the Orin's own 120-question S-EMBER sample, every row on the same questions —			
8B-I (Orin llama.cpp F16)	120	35.0	[27.1, 43.9]
8B-T (Orin llama.cpp F16)	120	36.7	[28.6, 45.6]
4B-I (A6000 llama.cpp F16)	120	31.7	[24.0, 40.4]
4B-T (A6000 llama.cpp F16)	120	36.7	[28.6, 45.6]
32B-I RICH K1 (H100 vLLM)	120	44.2	[35.6, 53.1]
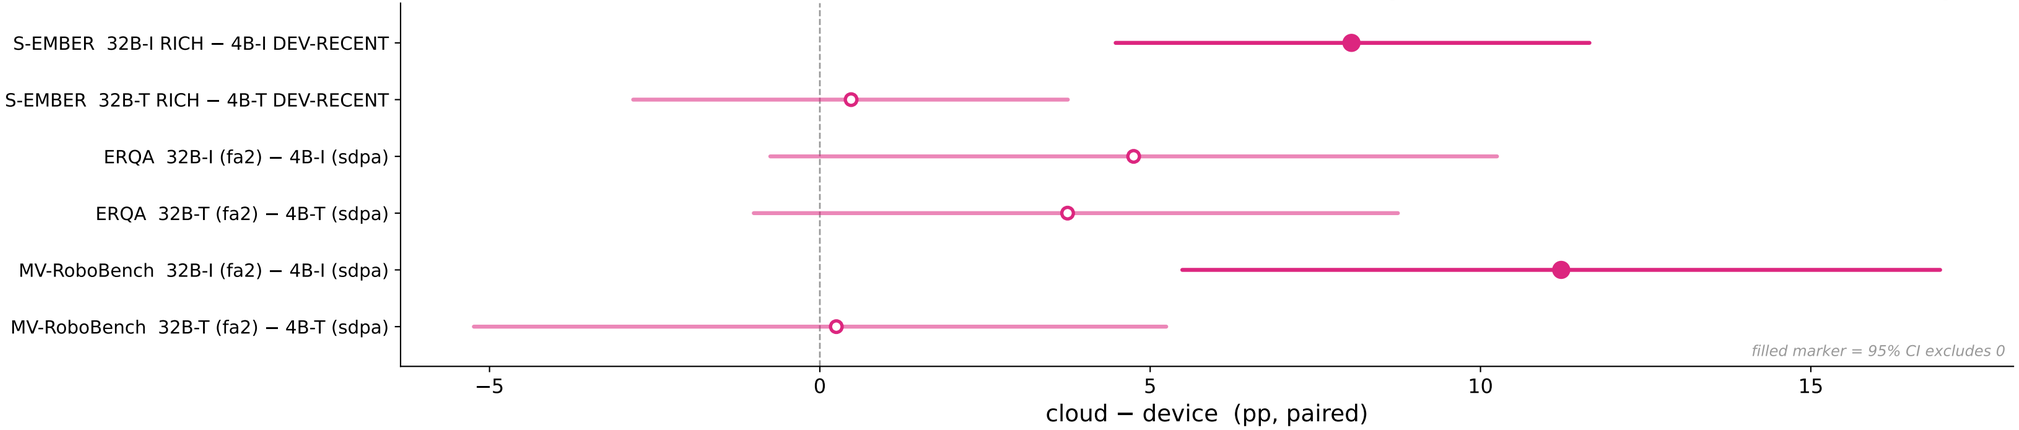
Paired differences, cloud minus device (S-EMBER: video-clustered bootstrap 2,000 / seed 0; ERQA & MV-RoboBench: paired question-level bootstrap 2,000 / seed 0)