What does local reasoning cost? Jetson AGX Orin, llama.cpp, DEV-RECENT 37-frame inputs (~12.4k tokens), 120-question sample (Q8_0: 40)

A. Time to first token — Jetson AGX Orin, llama.cpp
red lines: H100 32B vLLM at an 11k-token history, TTFT only, no network

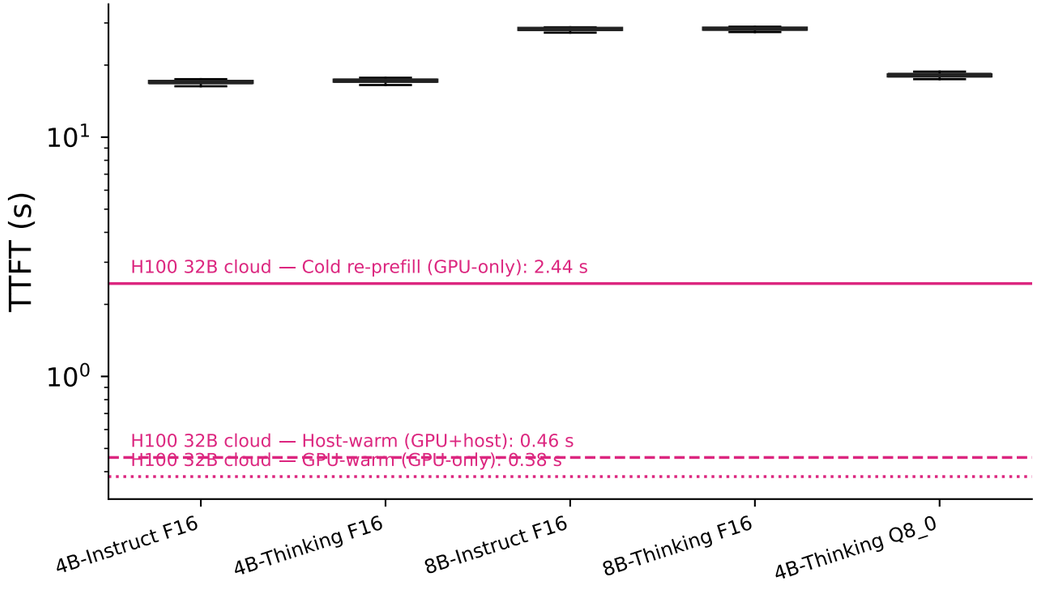

B. Full response latency (log) — Orin llama.cpp
Instruct is prefill-bound; Thinking adds a long, heavy tail

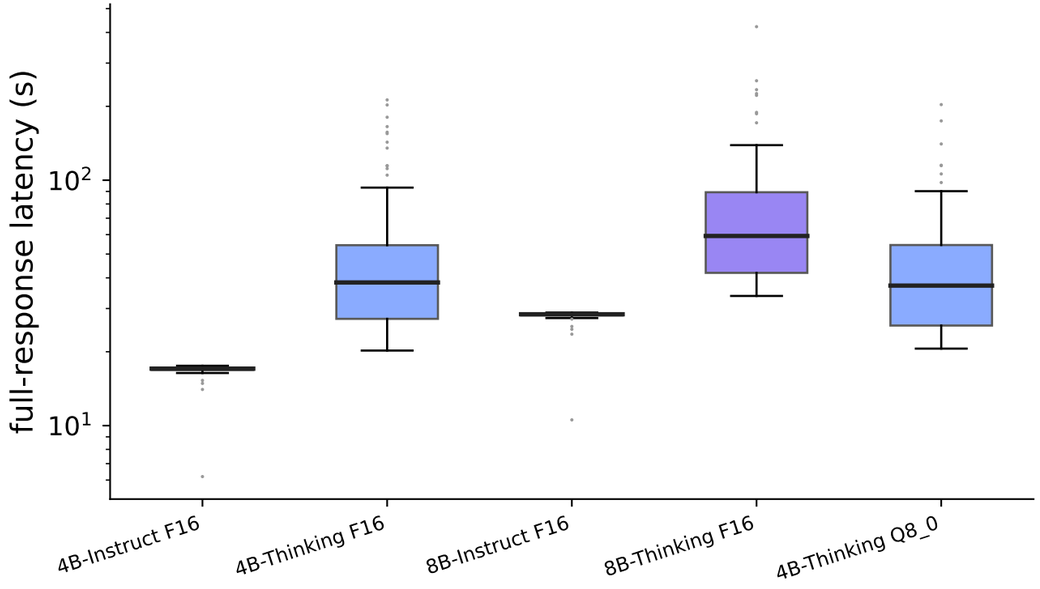

C. Reasoning-token distribution per Thinking model (Orin llama.cpp)

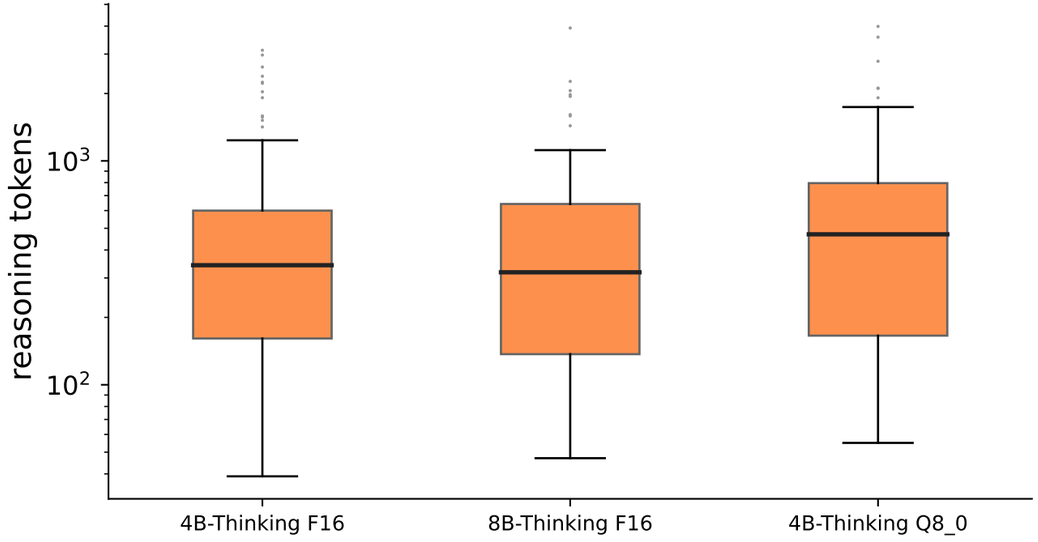
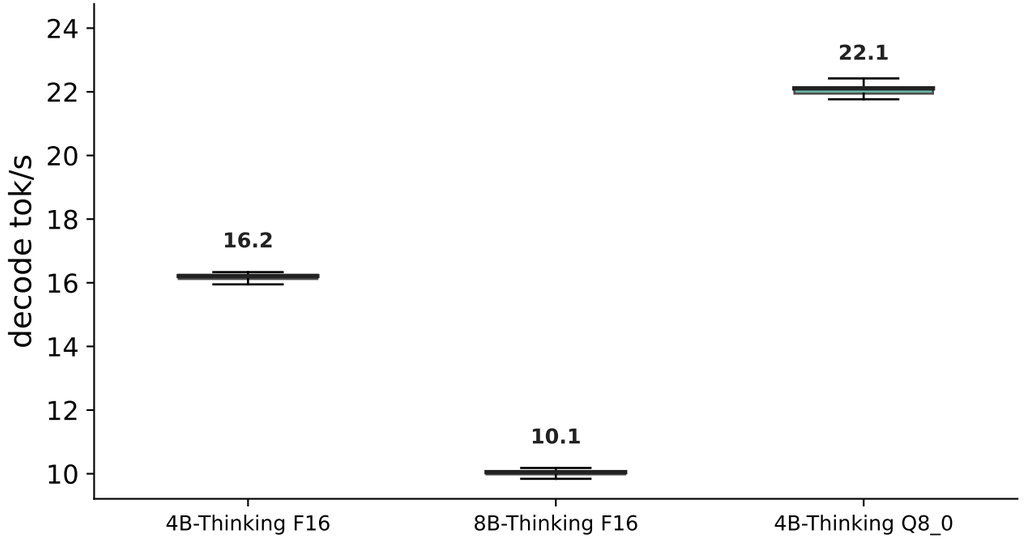
D. Decode throughput from the per-token timestamp arrays (tokens_*.jsonl.gz)