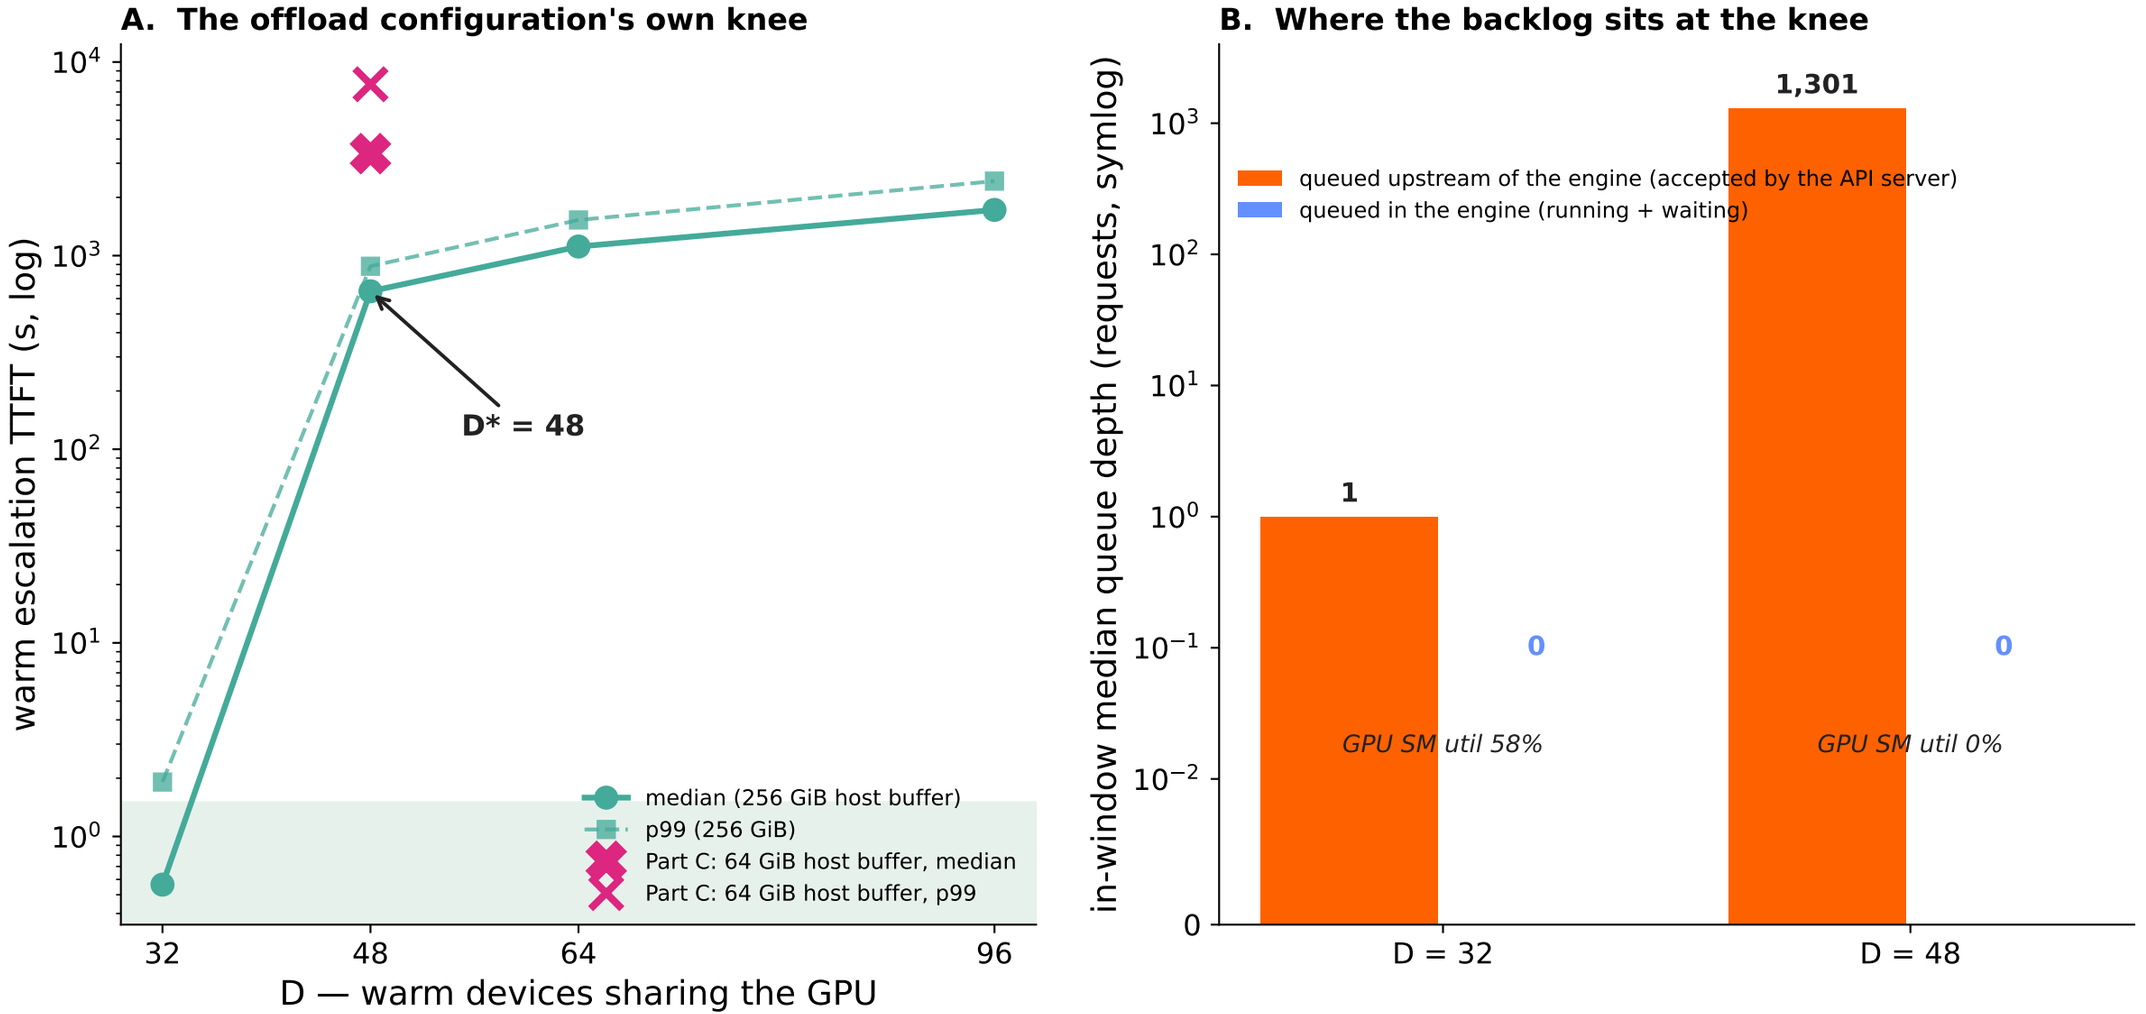Backup — StudyRH6: the knee of the GPU+host configuration is an API ingestion bottleneck (protocol), not GPU
H100 32B, vLLM, D2 fixed-stride histories; at D* = 48 the GPU is idle and the engine queue is empty while 1,301 requests wait in the API server