A larger cloud model is more accurate — and faster

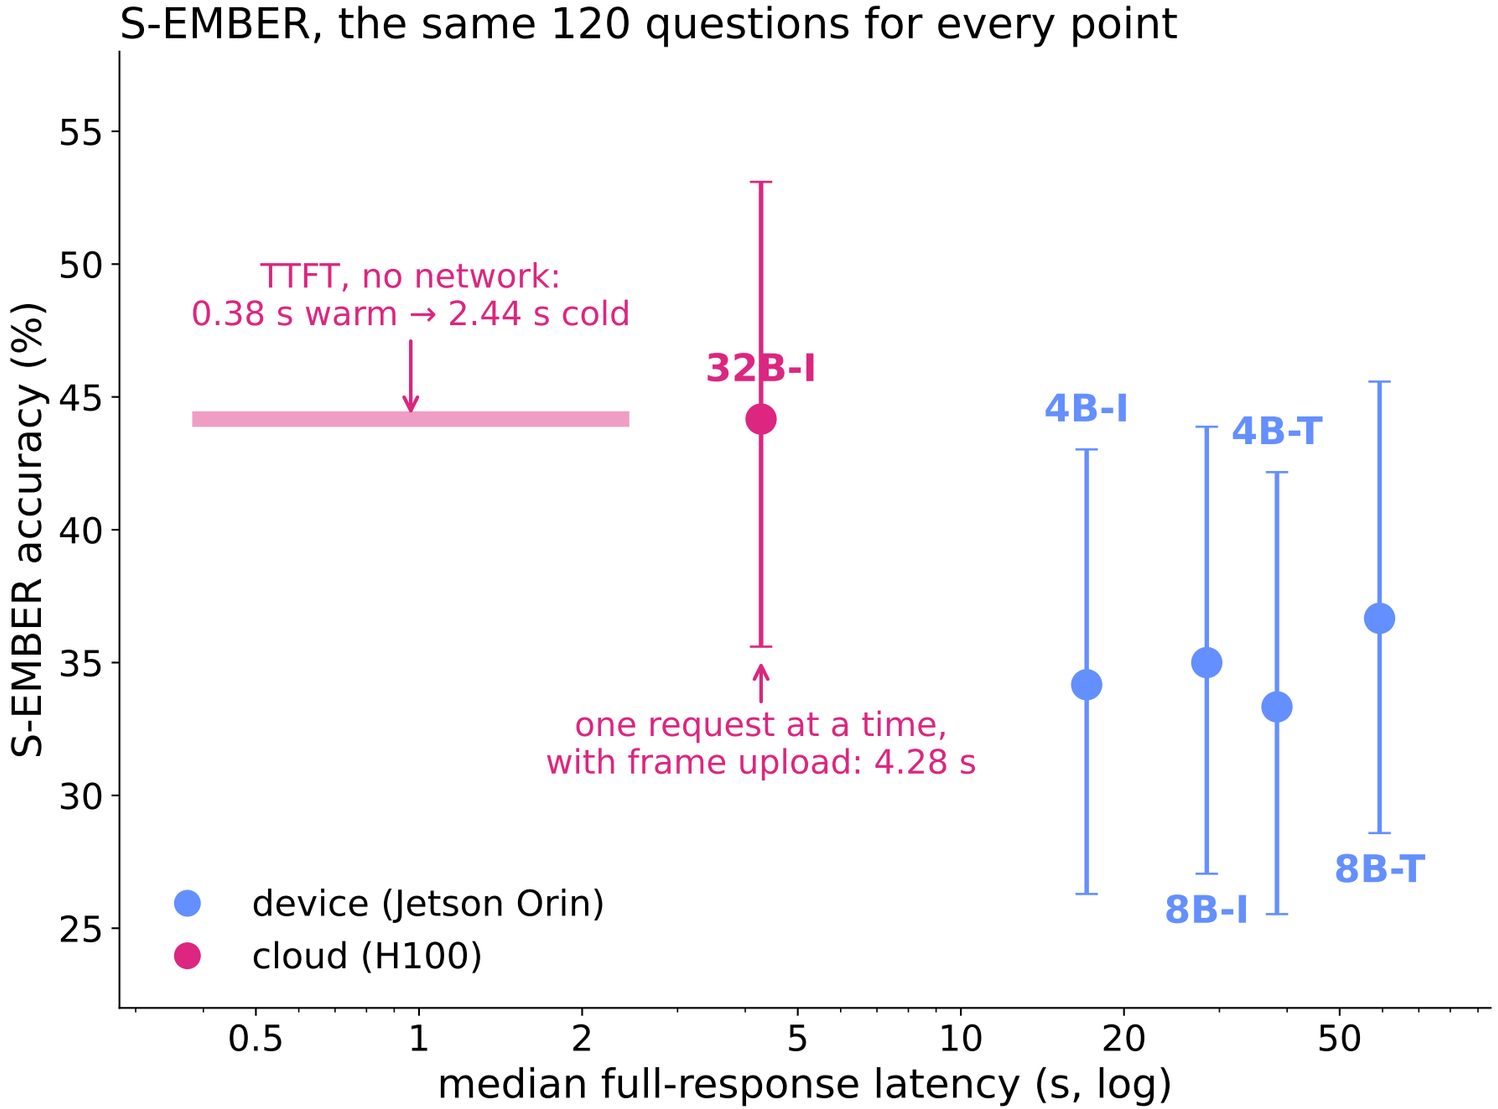

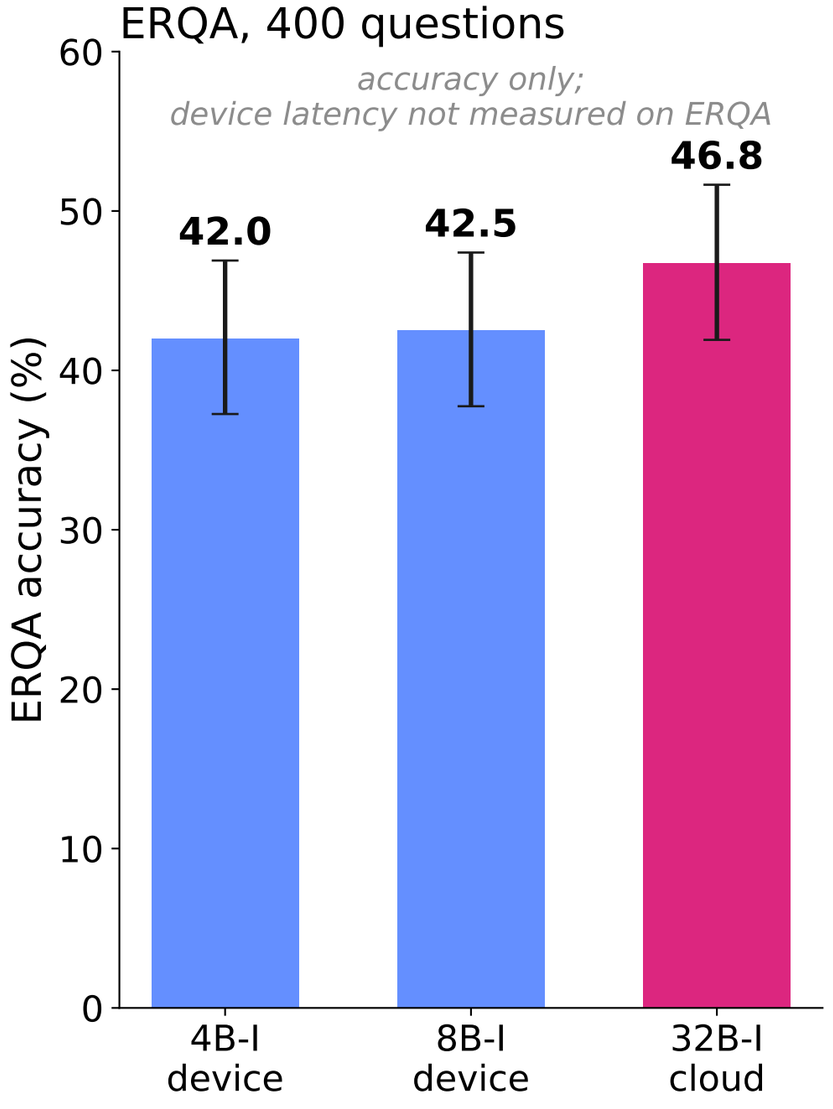

S-EMBER: device = Jetson AGX Orin, llama.cpp F16, DEV-RECENT 37-frame inputs, n=120; cloud = H100 NVL, Qwen3-VL-32B-Instruct, vLLM, RICH history (R1 K1), single-request latency n=50; TTFT bar = StudyRH5 Part A, 11k history, no network. ERQA: device 4B = 3090 Ti HF sdpa, device 8B = A6000 HF sdpa, cloud 32B = H100 HF fa2, n=400 (backend-mismatched). Error bars: Wilson 95%.