The cloud helps for two reasons: model size, and history the device has pruned

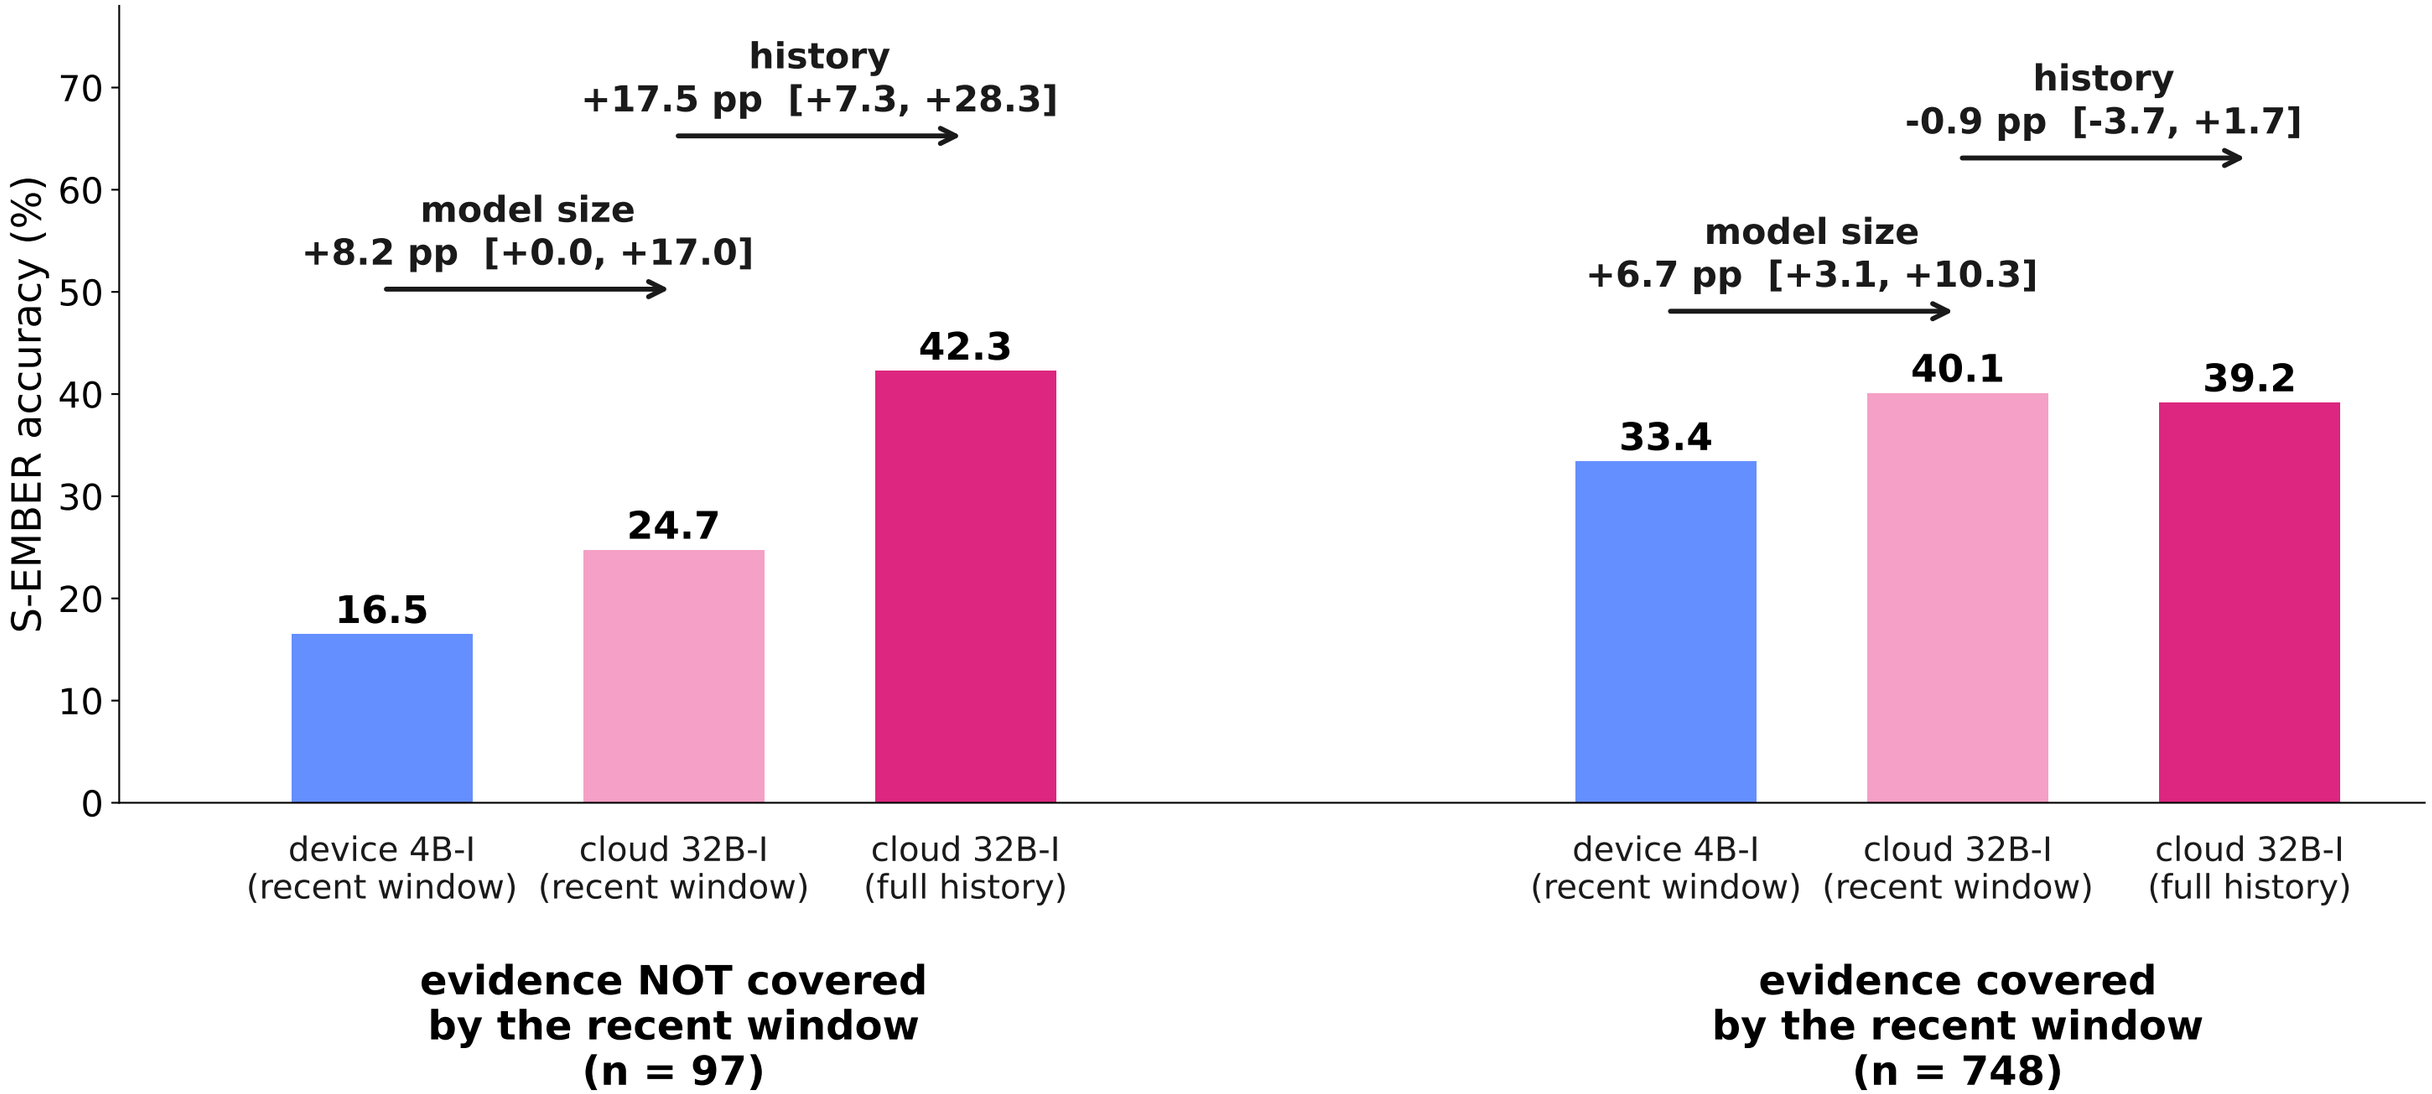

S-EMBER, 845 questions, split by the recent window's evidence coverage (study_r1 K3 `evidence_covered`). Device = A6000 llama.cpp F16, DEV-RECENT (accuracy_C1); cloud = H100 NVL, Qwen3-VL-32B-Instruct, vLLM, RECENT (K3) and RICH (K1). Device↔cloud arrow is backend-mismatched. CIs: video-clustered bootstrap, 2,000 resamples, seed 0.