# More thinking does not buy accuracy at either tier

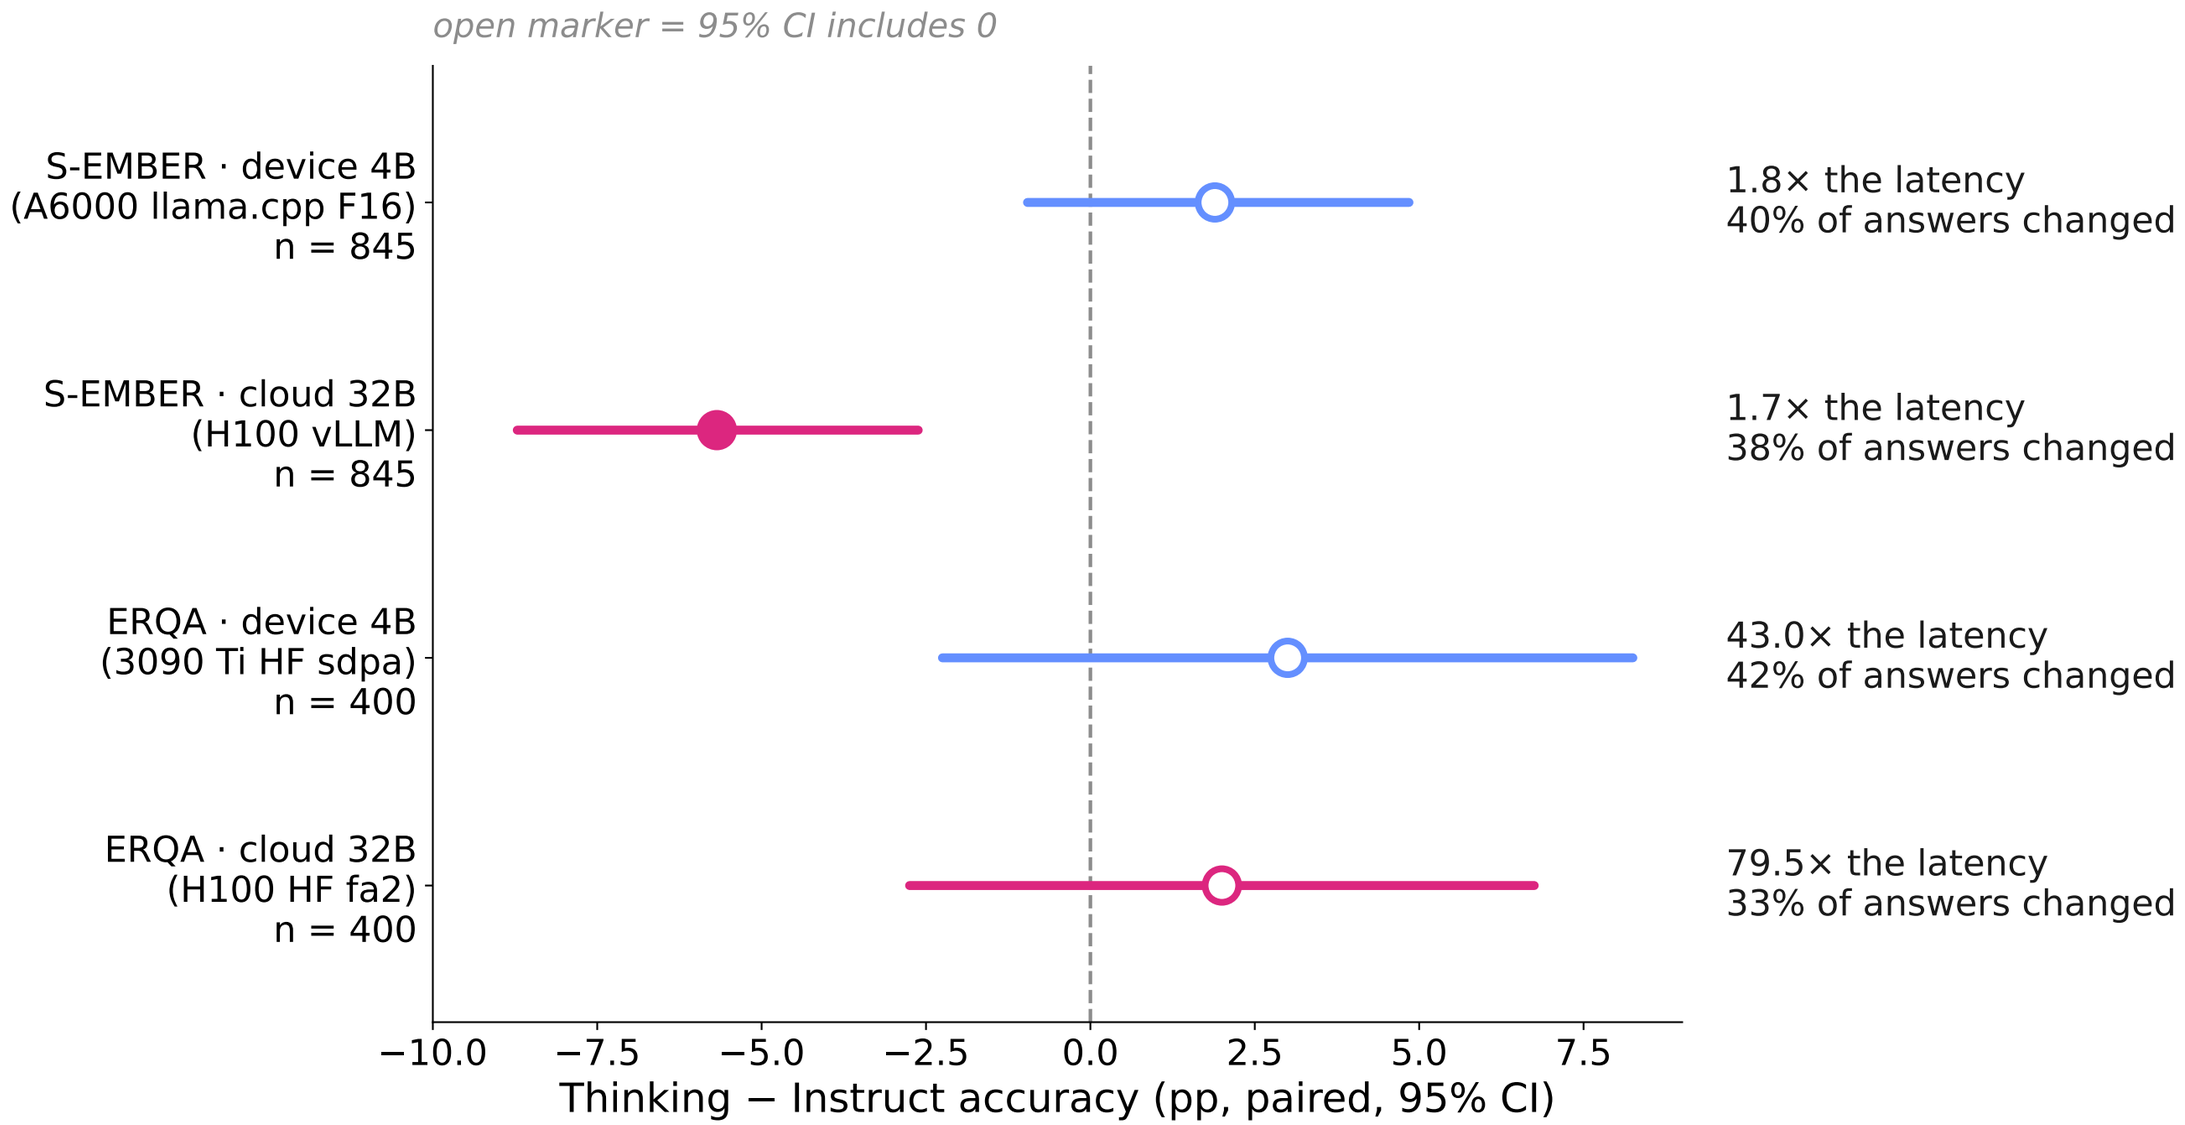


S-EMBER rows: video-clustered bootstrap, 2,000 resamples, seed 0 (device = A6000 llama.cpp F16 DEV-RECENT; cloud = H100 vLLM RICH). ERQA rows: paired question-level bootstrap, 2,000 resamples, seed 0 (device = 3090 Ti HF sdpa; cloud = H100 HF fa2). Latency multiples are medians on the same stack as the row.